

Exercises for "A Beginner's Introduction to Pydata: How to Build a Minimal Recommendation System"

Systems check

```
In [1]: import numpy as np
import pandas as pd
import tables as tb
!find ./data

./data
./data/.DS_Store
./data/ml-1m
./data/ml-1m/movies.dat
./data/ml-1m/ratings.dat
./data/ml-1m/README
./data/ml-1m/users.dat
./data/movielens_test.csv
./data/movielens_train.csv
```

How to load (a subset of) the MovieLens dataset

```
In [2]: # subset version (hosted notebook)
movielens_train = pd.read_csv('data/movielens_train.csv', index_col=0)
movielens_test = pd.read_csv('data/movielens_test.csv', index_col=0)
print movielens_train
print movielens_test

<class 'pandas.core.frame.DataFrame'>
Int64Index: 5838 entries, 593263 to 466639
Data columns:
user_id      5838  non-null values
movie_id     5838  non-null values
rating       5838  non-null values
timestamp    5838  non-null values
gender       5838  non-null values
age          5838  non-null values
occupation   5838  non-null values
zip          5838  non-null values
title        5838  non-null values
genres       5838  non-null values
for_testing  5838  non-null values
dtypes: bool(1), int64(6), object(4)
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2668 entries, 693323 to 713194
Data columns:
user_id      2668  non-null values
movie_id     2668  non-null values
rating       2668  non-null values
timestamp    2668  non-null values
gender       2668  non-null values
age          2668  non-null values
occupation   2668  non-null values
zip          2668  non-null values
title        2668  non-null values
genres       2668  non-null values
for_testing  2668  non-null values
dtypes: bool(1), int64(6), object(4)

```

Numpy Questions: Indexing

1. Access an individual element in a multidimensional array

```

In [3]: # given the following ndarray, access the its third element
arr = np.arange(10)
arr

```

```

Out[3]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])

```

2. Access the last column of of a 2d array

```

In [4]: # given the following ndarray, access its last column
arr = np.array([[5,4,2,5],[4,5,1,12],[0,1,5,4]])
arr

```

```

Out[4]: array([[ 5,  4,  2,  5],
               [ 4,  5,  1, 12],
               [ 0,  1,  5,  4]])

```

4. Select all elements from a 2d array that are larger than zero

```

In [5]: # given the following ndarray, obtain all elements that are larger than zero
arr = np.random.randn(5)
arr

```

```

Out[5]: array([-0.71949072,  0.58394499,  0.66661434,  0.75513099, -0.64750107])

```

```
Out[5]: array([ 0.71249072,  0.38324499,  0.00001434,  0.75519099, -0.04750107])
```

5. Set a portion of an array to the same scalar value

```
In [6]: # given the following ndarray, set the last two elements to 10  
arr = np.zeros(6)  
arr
```

```
Out[6]: array([ 0.,  0.,  0.,  0.,  0.,  0.])
```

Numpy Questions: Operations

1. Compute the sum of a 1D array

```
In [7]: # given the following ndarray, compute its sum  
arr = np.random.randn(5)  
arr
```

```
Out[7]: array([-0.25262032,  2.54477542, -1.39308539,  0.8186857 ,  0.30090793])
```

2. Compute the mean of a 1D array

```
In [8]: # given the following ndarray, compute its mean  
arr = np.random.randn(5)  
arr
```

```
Out[8]: array([ 0.56138573,  0.95797777, -1.24147562, -0.08016767,  1.55450634])
```

3. How do you detect the presence of NaNs in an array?

```
In [9]: # given the following ndarray, detect all elements that are nans  
arr = np.array([np.nan] * 10)  
arr[2:4] = 5  
arr
```

```
Out[9]: array([ nan,  nan,   5.,   5.,  nan,  nan,  nan,  nan,  nan,  nan])
```

Pandas questions: Series and DataFrames

1. Adding and deleting a column in a dataframe

```
In [10]: # given the following DataFrame, add a new column to it  
df = pd.DataFrame({'col1': [1,2,3,4]})  
df
```

Out[10]:

	col1
0	1
1	2
2	3
3	4

2. Adding and deleting a row in a dataframe

```
In [11]: # given the following DataFrame, delete row 'd' from it  
df = pd.DataFrame({'col1': [1,2,3,4]}, index = ['a','b','c','d'])  
df
```

Out[11]:

	col1
a	1
b	2
c	3
d	4

3. Creating a DataFrame from a few Series

```
In [12]: # given the following three Series, create a DataFrame such that it holds  
ser_1 = pd.Series(np.random.randn(6))  
ser_2 = pd.Series(np.random.randn(6))  
ser_3 = pd.Series(np.random.randn(6))
```

Pandas questions: indexing

1. Indexing into a specific column

```
In [32]: # given the dataframe 'movielens' that we loaded in the previous step, try
         # into the 'zip' column
         movielens[?]
```

2. Label-based indexing

```
In [33]: # using the same 'movielens' dataframe, index into the row whose index is
         movielens.ix[?]
```

Reco systems questions: estimation functions

1. Simple content filtering using mean ratings

```
In [ ]: # write an 'estimate' function that computes the mean rating of a particular
         def estimate(user_id, movie_id):
             # first, index into all ratings by this user
             # second, compute the mean of those ratings
             # return

         # try it out for a user_id, movie_id pair
         estimate(4653, 2648)
```

2. Simple collaborative filtering using mean ratings

```
In [ ]: # write an 'estimate' function that computes the mean rating of a particular
         def estimate(user_id, movie_id):
             # first, index into all ratings of this movie
             # second, compute the mean of those ratings
```

```
# return

# try it out for a user_id, movie_id pair
estimate(4653, 2648)
```

Mini-Challenge

These are the two functions that you will need to test your estimate method.

```
In [13]: def compute_rmse(y_pred, y_true):
          """ Compute Root Mean Squared Error. """
          return np.sqrt(np.mean(np.power(y_pred - y_true, 2)))
```

```
In [14]: def evaluate(estimate_f):
          """ RMSE-based predictive performance evaluation with pandas. """
          ids_to_estimate = zip(movielens_test.user_id, movielens_test.movie_id)
          estimated = np.array([estimate_f(u,i) for (u,i) in ids_to_estimate])
          real = movielens_test.rating.values
          return compute_rmse(estimated, real)
```

With those, you can test for performance with the following line, which assumes that your function is called `my_estimate_func`:

```
In [ ]: print 'RMSE for my estimate function: %s' % evaluate(my_estimate_func)
```

Pytables questions: file and node creation

1. Create a PyTables file in your working environment

```
In [ ]: # write your answer in this code block
```

2. Within the file you created, create a new group

```
In [ ]: # write your answer in this code block
```

3. Within the group you created, create a new array of integers

and save it

```
In [ ]: # write your answer in this code block
```

4. For the group created, set a datetime attribute, with the value of 'utcnow'

```
In [ ]: # write your answer in this code block
```